

Final_Project

August 16, 2025

```
[1]: import numpy as np
import pandas as pd
from math import lgamma, log
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.optimize import minimize
from pathlib import Path

sns.set_theme(style="whitegrid", rc={"figure.dpi": 120})

# Auto-root (specifying path)
ROOT = Path(__file__).resolve().parent if "__file__" in globals() else Path.
    ↪cwd().resolve()
DATA_DIR = ROOT / "data"
POINTS_DIR = ROOT / "points"

training_csvs = [DATA_DIR / "E0_19_20.csv", DATA_DIR / "E0_20_21.csv", DATA_DIR /
    ↪ "E0_21_22.csv"]
backtest_fixtures_csvs = [DATA_DIR / "epl_22_23_fixtures.csv", DATA_DIR /
    ↪ "epl_23_24_fixtures.csv"]
backtest_actual_pts_csvs = [POINTS_DIR / "epl_2022_23_points.csv", POINTS_DIR /
    ↪ "epl_2023_24_points.csv"]
future_fixtures_csvs = [DATA_DIR / "epl_24_25_fixtures.csv", DATA_DIR /
    ↪ "epl_25_26_fixtures.csv"]

# Monte Carlo simulations per season
N_SIMS = 500

# 1) Data Loading & Engineering
def load_and_engineer_data(path: str) -> pd.DataFrame:
    cols = ["Date", "HomeTeam", "AwayTeam", "FTHG", "FTAG"]
    df = pd.read_csv(path, usecols=cols, dayfirst=True, parse_dates=["Date"])

    # Basic NA handling
```

```

df = df.dropna(how="any", subset=["Date", "HomeTeam", "AwayTeam", "FTHG",
↪ "FTAG"])

# Strip extra spaces from team names
df["HomeTeam"] = df["HomeTeam"].str.strip()
df["AwayTeam"] = df["AwayTeam"].str.strip()

# Ensure goals are integers and non-negative
df["FTHG"] = pd.to_numeric(df["FTHG"], errors="coerce").fillna(0).
↪ astype(int)
df["FTAG"] = pd.to_numeric(df["FTAG"], errors="coerce").fillna(0).
↪ astype(int)
df = df[(df["FTHG"] >= 0) & (df["FTAG"] >= 0)]

# Sort by date
df = df.sort_values("Date").reset_index(drop=True)

# EDA: Print missing values summary
missing_summary = df.isnull().sum()
if missing_summary.any():
    print(f"[WARNING] Missing values found in {path.name}:")
    print(missing_summary[missing_summary > 0])
else:
    print(f"[INFORMATION] No missing values in {path.name}")

return df

def get_unique_teams(df: pd.DataFrame):
    return sorted(pd.unique(pd.concat([df["HomeTeam"], df["AwayTeam"]])))

```

[]: # 2) Elo Ratings with Decay

```

def compute_elo(df: pd.DataFrame, teams, K: float = 30.0, decay: float = 1.0):
    """
    Compute Elo with decay. Effective step = (1 - decay) * K.
    Returns (df_with_elos, final_elo_dict).
    """
    elo = {t: 1500.0 for t in teams}
    out = df.copy()
    out["HomeElo"] = np.nan
    out["AwayElo"] = np.nan

    for i, r in out.iterrows():
        h, a = r["HomeTeam"], r["AwayTeam"]
        Rh, Ra = elo[h], elo[a]
        out.at[i, "HomeElo"] = Rh
        out.at[i, "AwayElo"] = Ra

```

```

    Eh = 1.0 / (1.0 + 10.0 ** ((Ra - Rh) / 400.0))
    Ea = 1.0 - Eh

    if r["FTHG"] > r["FTAG"]:
        Sh = 1.0
    elif r["FTHG"] < r["FTAG"]:
        Sh = 0.0
    else:
        Sh = 0.5
    Sa = 1.0 - Sh

    # Simplified & correct with decay factor
    elo[h] = Rh + (1.0 - decay) * K * (Sh - Eh)
    elo[a] = Ra + (1.0 - decay) * K * (Sa - Ea)

    return out, elo

```

```

[3]: # 3) Bivariate Poisson pmf & Likelihood
def _log_dpois(k: int, lam: float) -> float:
    if lam <= 0:
        return -np.inf if k > 0 else 0.0
    return k * log(lam) - lam - lgamma(k + 1.0)

def _logsumexp(a: np.ndarray) -> float:
    m = np.max(a)
    return m + np.log(np.sum(np.exp(a - m)))

def log_bipois(x: int, y: int, lam1: float, lam2: float, lam3: float) -> float:
    kmax = min(x, y)
    logs = []
    for k in range(kmax + 1):
        logs.append(_log_dpois(x - k, lam1) + _log_dpois(y - k, lam2) +
↪ _log_dpois(k, lam3))
    return _logsumexp(np.array(logs))

def bipois_nll(par: np.ndarray, df_e: pd.DataFrame, teams, gamma: float,
↪ lambda_ridge: float) -> float:
    """Negative penalized log-likelihood."""
    n = len(teams)
    atk = par[0:n].copy()
    dfn = par[n:2*n].copy()
    hAdv = par[2*n:3*n].copy()
    lam3 = np.exp(par[3*n])
    idx = {t: i for i, t in enumerate(teams)}

    ll = -lambda_ridge * (np.sum(atk**2) + np.sum(dfn**2))

```

```

for _, r in df_e.iterrows():
    iH = idx[r["HomeTeam"]]; iA = idx[r["AwayTeam"]]
    ratio = (r["HomeElo"] / r["AwayElo"]) if r["AwayElo"] > 0 else 1.0
    lam1 = np.exp(atk[iH] - dfn[iA] + hAdv[iH]) * (ratio ** gamma)
    lam2 = np.exp(atk[iA] - dfn[iH]) / (ratio ** gamma)
    ll += log_bipois(int(r["FTHG"]), int(r["FTAG"]), float(lam1),
    ↪ float(lam2), float(lam3))
    return -ll # minimize

def fit_bipois(df_e: pd.DataFrame, teams, gamma: float = 0.1, lambda_ridge:
    ↪ float = 0.01):
    """BFGS fit with zero-mean constraints on attack & defence."""
    n = len(teams)
    init = np.zeros(3*n + 1, dtype=float)
    init[3*n] = np.log(0.1) # initial shared component

    def obj(p):
        p = p.copy()
        p[0:n] -= np.mean(p[0:n])
        p[n:2*n] -= np.mean(p[n:2*n])
        return bipois_nll(p, df_e, teams, gamma, lambda_ridge)

    res = minimize(obj, init, method="BFGS", options={"maxiter": 400, "gtol":
    ↪ 1e-5})
    p = res.x.copy()
    p[0:n] -= np.mean(p[0:n])
    p[n:2*n] -= np.mean(p[n:2*n])
    idx = {t: i for i, t in enumerate(teams)}
    return {"par": p, "idx": idx}

```

```

[4]: # 4) Simulations
_rng = np.random.default_rng(1)

def simulate_match(h: str, a: str, par: np.ndarray, idx: dict, elo: dict, gamma:
    ↪ float):
    n = len(idx)
    i = idx[h]; j = idx[a]
    atk = par[0:n]; dfn = par[n:2*n]; hAdv = par[2*n:3*n]; lam3 = np.
    ↪ exp(par[3*n])

    ratio = (elo[h] / elo[a]) if elo[a] > 0 else 1.0
    lam1 = np.exp(atk[i] - dfn[j] + hAdv[i]) * (ratio ** gamma)
    lam2 = np.exp(atk[j] - dfn[i]) / (ratio ** gamma)

    k = _rng.poisson(lam3)
    x = _rng.poisson(lam1)
    y = _rng.poisson(lam2)

```

```

        return x + k, y + k

def simulate_season(fixture: pd.DataFrame, par: np.ndarray, idx: dict, elo: dict,
    gamma: float):
    pts = {team: 0 for team in idx.keys()}
    for _, r in fixture.iterrows():
        h, a = r["HomeTeam"], r["AwayTeam"]
        x, y = simulate_match(h, a, par, idx, elo, gamma)
        if x > y:
            pts[h] += 3
        elif x < y:
            pts[a] += 3
        else:
            pts[h] += 1; pts[a] += 1
    return pts

def monte_carlo(fixture: pd.DataFrame, par: np.ndarray, idx: dict, elo: dict,
    gamma: float, n: int = 100):
    teams = list(idx.keys())
    sims = []
    for _ in range(n):
        pts = simulate_season(fixture, par, idx, elo, gamma)
        sims.append([pts[t] for t in teams])
    mat = np.array(sims).T # teams x n
    return pd.DataFrame(mat, index=teams)

```

```

[ ]: # 5) Hyper-parameter Search (multiple backtests)
def search_params(train_df: pd.DataFrame, fixture_csvs, actual_csvs,
    teams_all):
    fixture_list = [pd.read_csv(f)[["HomeTeam", "AwayTeam"]] for f in
    fixture_csvs]
    actuals_list = []
    for f in actual_csvs:
        ap = pd.read_csv(f).rename(columns={"Points": "ActualPts"})
        actuals_list.append(ap[["Team", "ActualPts"]])

    best = {"MAE": np.inf}
    for K in [20,30,40]: #k grid (keeping it small for speed, can
    expand if needed)
        df_e, elo_vec = compute_elo(train_df, teams_all, K=K, decay=0.995)
        for g in [0.04,0.06]: # gamma grid (keeping it small for speed, can
        expand if needed)
            for lam in [0.02]: # lambda grid (keeping it small for speed, can
            expand if needed)
                maes = []
                for fx, ap in zip(fixture_list, actuals_list):

```

```

        fit = fit_bipois(df_e, teams_all, gamma=g, lambda_ridge=lam)
        sim_mat = monte_carlo(fx, fit["par"], fit["idx"], elo_vec,
↪g, n=N_SIMS)

        pred = sim_mat.mean(axis=1).rename("Pred").reset_index()
        pred = pred.rename(columns={"index": "Team"})
        cmp = pd.merge(pred, ap, on="Team", how="inner")
        maes.append(np.mean(np.abs(cmp["Pred"] - cmp["ActualPts"])))
    mae_all = float(np.mean(maes))
    if mae_all < best["MAE"]:
        best = {"MAE": mae_all, "K": K, "gamma": g, "lambda": lam}

    return best

```

```

[6]: # Helpers: plots & finish-position probabilities
def plot_points_boxplot(sim_mat: pd.DataFrame, team_order, title="Simulated_
↪Points Distribution"):
    long = sim_mat.loc[team_order].T.melt(var_name="Team", value_name="Points")
    plt.figure(figsize=(5,7))
    sns.boxplot(data=long, x="Points", y="Team", fliersize=1)
    plt.title(title); plt.xlabel("Points"); plt.ylabel("")
    plt.tight_layout(); plt.show()

def compute_finish_probs(sim_mat: pd.DataFrame, team_order):
    """
    Returns:
        prob_df: DataFrame [teams x positions] with P(position==p)
        rank_long: tidy long-form for heatmap
    """
    sim_sub = sim_mat.loc[team_order].values # T x N
    T, N = sim_sub.shape
    ranks = np.zeros_like(sim_sub, dtype=int)
    rng = np.random.default_rng(0)
    for j in range(N):
        pts = sim_sub[:, j].astype(float) + 1e-8 * rng.uniform(size=T)
        order = np.argsort(-pts) # high to low
        r = np.empty_like(order)
        r[order] = np.arange(1, T + 1)
        ranks[:, j] = r

    probs = np.zeros((T, T))
    for p in range(1, T + 1):
        probs[:, p - 1] = np.mean(ranks == p, axis=1)

    prob_df = pd.DataFrame(probs, index=team_order, columns=[str(i) for i in
↪range(1, T + 1)])
    rank_long = prob_df.reset_index().melt(id_vars="index",
↪var_name="Position", value_name="Prob")
    rank_long = rank_long.rename(columns={"index": "Team"})

```

```

    # order: positions 1..N left ->right; best teams at TOP
    rank_long["Team"] = pd.Categorical(rank_long["Team"],
    ↪categories=list(reversed(team_order)), ordered=True)
    rank_long["Position"] = pd.Categorical(rank_long["Position"],
    ↪categories=[str(i) for i in range(1, T + 1)], ordered=True)
    return prob_df, rank_long

def plot_finish_heatmap(rank_long, title="Finish Position Probability"):
    pivot = rank_long.pivot_table(
        index="Team",
        columns="Position",
        values="Prob",
        aggfunc="first",
        observed=False
    ).fillna(0)

    plt.figure(figsize=(8,6))
    ax = sns.heatmap(pivot, cmap="Blues", vmin=0, vmax=1,
                     cbar_kws={"label": "Probability (dark = higher)"})
    ax.set_title(title)
    ax.set_xlabel("League Position"); ax.set_ylabel("")
    plt.tight_layout(); plt.show()

def plot_outcome_bars(prob_df: pd.DataFrame, team_order, title="Key Outcome
    ↪Probabilities"):
    n_teams = prob_df.shape[0]
    title_prob = prob_df["1"].values if "1" in prob_df.columns else np.
    ↪zeros(n_teams)
    top4_prob = prob_df[[str(i) for i in range(1, min(4, n_teams) + 1)]].
    ↪sum(axis=1).values
    releg_prob = prob_df[[str(i) for i in range(n_teams - 2, n_teams + 1)]].
    ↪sum(axis=1).values

    bucket = pd.DataFrame({"Team": team_order, "Title": title_prob, "Top-4":
    ↪top4_prob, "Relegation": releg_prob})
    long = bucket.melt(id_vars="Team", var_name="Outcome",
    ↪value_name="Probability")
    long["Team"] = pd.Categorical(long["Team"], categories=team_order,
    ↪ordered=True)

    plt.figure(figsize=(7,7))
    sns.barplot(data=long, x="Probability", y="Team", hue="Outcome", orient="h")
    plt.title(title); plt.xlabel("Probability"); plt.ylabel("")
    plt.legend(loc="upper right")
    plt.tight_layout(); plt.show()

```

```

[7]: # 6) Main: train, tune, fit, simulate future
def main():
    # 6.1 Build training set
    train_list = [load_and_engineer_data(p) for p in training_csvs]
    train_df = pd.concat(train_list, ignore_index=True)

    # Ensure every club in any fixture is included
    all_fixtures = pd.concat(
        [pd.read_csv(f)[["HomeTeam", "AwayTeam"]] for f in
        ↪(backtest_fixtures_csvs + future_fixtures_csvs)],
        ignore_index=True
    )
    teams_all = sorted(set(get_unique_teams(train_df)) |
                        set(all_fixtures["HomeTeam"]) |
                        set(all_fixtures["AwayTeam"]))

    # 6.2 Tune via backtests (single grid for speed)
    best = search_params(train_df, backtest_fixtures_csvs,
    ↪backtest_actual_pts_csvs, teams_all)
    best_K, best_gamma, best_lambda = best["K"], best["gamma"], best["lambda"]
    print(f"Tuned: K={best_K:.1f}, Gamma={best_gamma:.2f}, Lambda={best_lambda:.
    ↪3f} (MAE={best['MAE']:.2f})")

    # 6.3 Fit final model
    df_e, elo_vec = compute_elo(train_df, teams_all, K=best_K, decay=0.995)
    fit_res = fit_bipois(df_e, teams_all, gamma=best_gamma,
    ↪lambda_ridge=best_lambda)

    # 6.4 Simulate each future season + plots
    for csv in future_fixtures_csvs:
        csv = Path(csv) # ensure Path object
        label = csv.name # taking just name of the file like
        ↪"epl_25_26_fixtures.csv" as label

        fx = pd.read_csv(csv)[["HomeTeam", "AwayTeam"]]
        sim_mat = monte_carlo(fx, fit_res["par"], fit_res["idx"], elo_vec,
        ↪best_gamma, n=N_SIMS)
        med_pts = sim_mat.median(axis=1)
        curr20 = pd.unique(pd.concat([fx["HomeTeam"], fx["AwayTeam"]]))
        out = (
            pd.DataFrame({"Team": med_pts.index, "PredictedPts": med_pts.values})
            .query("Team in @curr20")
            .sort_values("PredictedPts", ascending=False)
            .reset_index(drop=True)
        )

```



```

print(f"\n=== Predictions for {label} ===")
print(out.to_string(index=False))

team_order = out["Team"].tolist()
plot_points_boxplot(sim_mat, team_order, title=f"Simulated Points_
↳Distribution ({label})")
prob_df, rank_long = compute_finish_probs(sim_mat, team_order)
plot_finish_heatmap(rank_long, title=f"Finish Position Probability_
↳({label})")
plot_outcome_bars(prob_df, team_order, title=f"Key Outcome_
↳Probabilities ({label})")
if __name__ == "__main__":
    main()


```

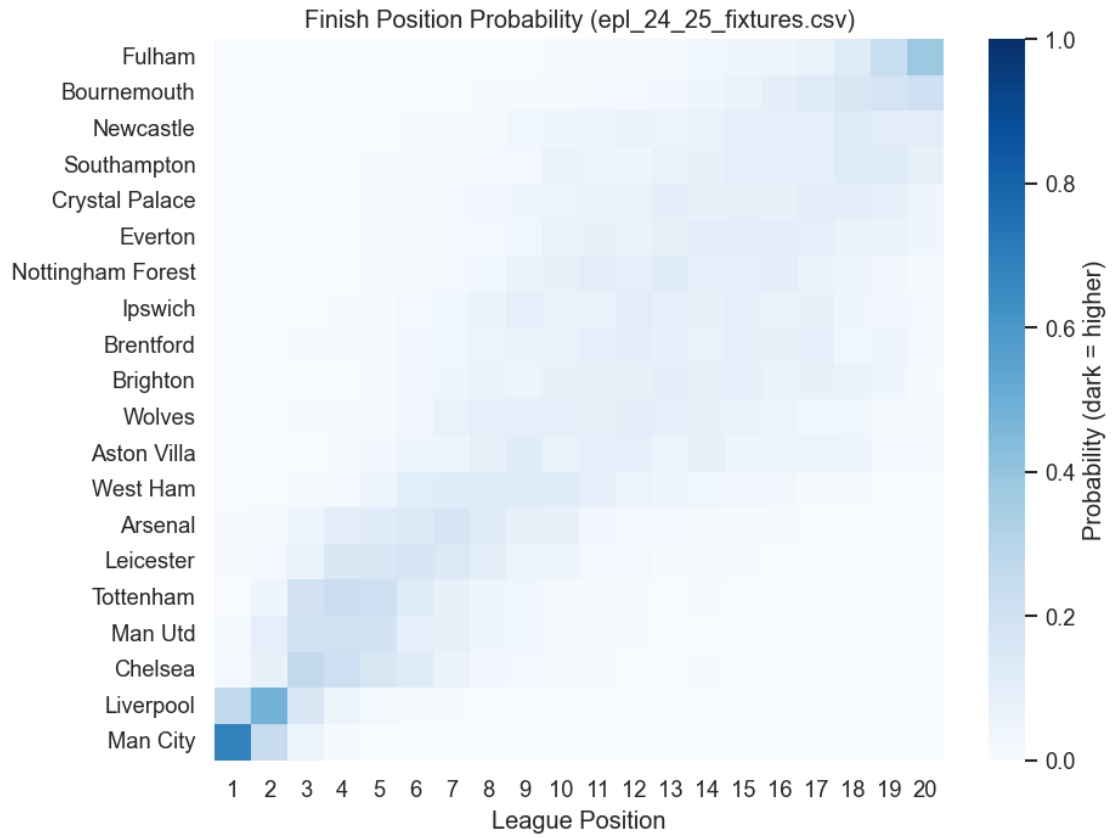
[INFORMATION] No missing values in E0_19_20.csv
[INFORMATION] No missing values in E0_20_21.csv
[INFORMATION] No missing values in E0_21_22.csv
Tuned: K=20.0, Gamma=0.06, Lambda=0.020 (MAE=11.67)

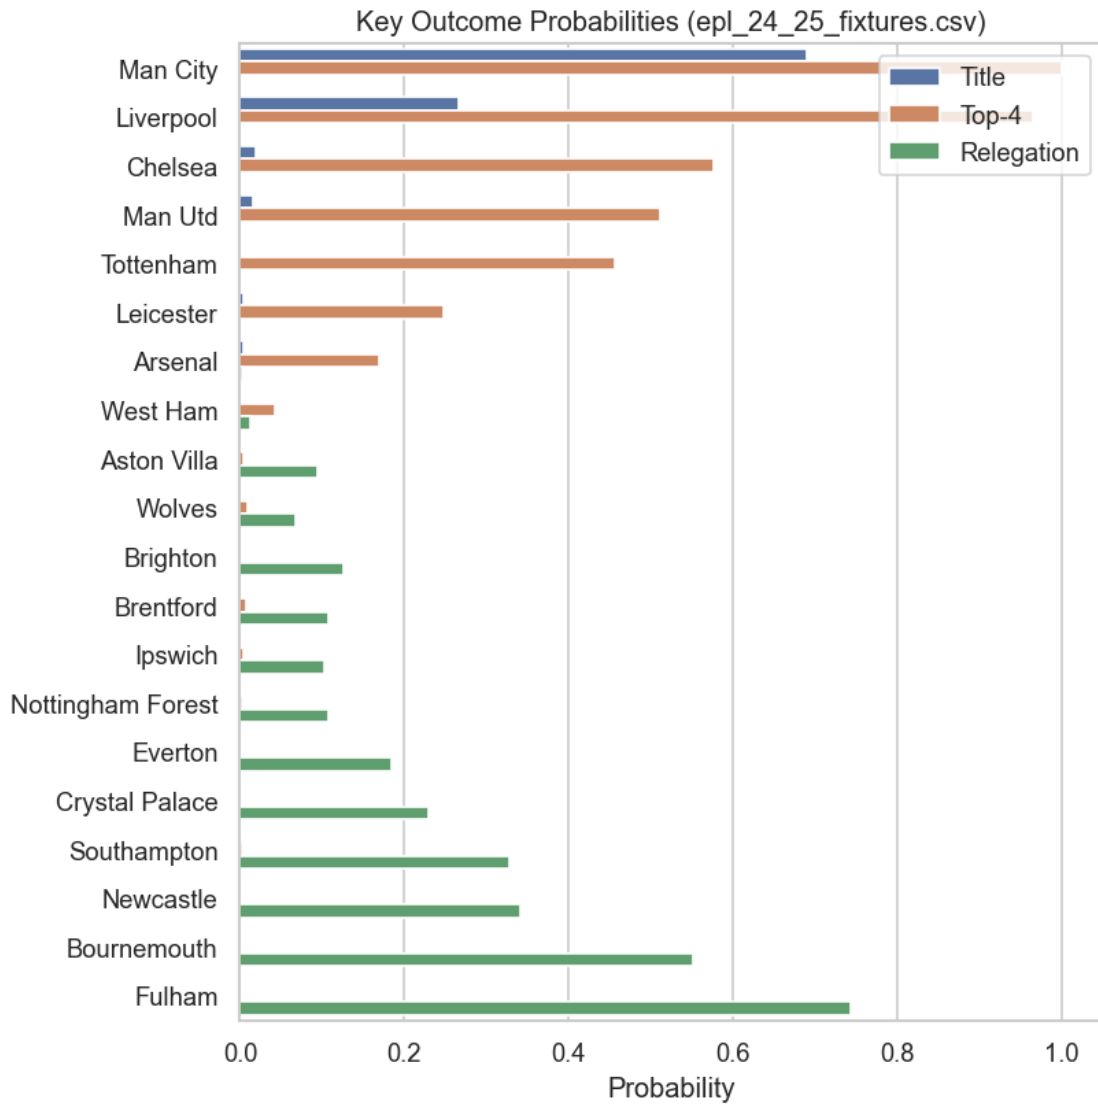
=== Predictions for epl_24_25_fixtures.csv ===

Team	PredictedPts
Man City	86.0
Liverpool	80.0
Chelsea	67.0
Man Utd	66.0
Tottenham	64.0
Leicester	60.0
Arsenal	58.0
West Ham	53.0
Aston Villa	47.0
Wolves	47.0
Brighton	45.0
Brentford	45.0
Ipswich	45.0
Nottingham Forest	44.0
Everton	43.0
Crystal Palace	42.0
Southampton	40.0
Newcastle	39.0
Bournemouth	35.0
Fulham	32.0

Simulated Points Distribution (epl_24_25_fixtures.csv)







=== Predictions for epl_25_26_fixtures.csv ===

Team	PredictedPts
Man City	87.0
Liverpool	80.0
Chelsea	68.0
Man Utd	68.0
Tottenham	64.0
Arsenal	59.0
West Ham	54.0
Wolves	49.0
Aston Villa	48.0
Brentford	46.0

Brighton	46.0
Sunderland	45.0
Nottingham Forest	44.5
Everton	43.0
Crystal Palace	43.0
Leeds	43.0
Burnley	41.0
Newcastle	40.0
Bournemouth	36.0
Fulham	33.0

Simulated Points Distribution (epl_25_26_fixtures.csv)

